

```
In [1]: # This Python 3 environment comes with many helpful analytics libraries installed
# It is defined by the kaggle/python Docker image: https://github.com/kaggle
# For example, here's several helpful packages to load
import seaborn as sns
import plotly.express as px
import matplotlib.pyplot as plt
import plotly.graph_objects as go
from tqdm import tqdm_notebook
import plotly.figure_factory as ff

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

import warnings
warnings.filterwarnings('ignore')

plt.style.use('fivethirtyeight')
%matplotlib inline
# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking run or pressing Shift+Enter) will list all files in the input directory

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# You can write up to 20GB to the current directory (/kaggle/working/) that
# You can also write temporary files to /kaggle/temp/, but they won't be saved
```

/kaggle/input/water-potability/water_potability.csv

```
In [2]: data=pd.read_csv('../input/water-potability/water_potability.csv')
data.head()
```

```
Out[2]:
```

	ph	Hardness	Solids	Chloramines	Sulfate	Conductivity	Organic_c
0	NaN	204.890455	20791.318981	7.300212	368.516441	564.308654	10.3
1	3.716080	129.422921	18630.057858	6.635246	NaN	592.885359	15.1
2	8.099124	224.236259	19909.541732	9.275884	NaN	418.606213	16.8
3	8.316766	214.373394	22018.417441	8.059332	356.886136	363.266516	18.4
4	9.092223	181.101509	17978.986339	6.546600	310.135738	398.410813	11.5

EDA

- ph-> pH of water
- Hardness-> Capacity of water to precipitate soap in mg/L
- Solids-> Total dissolved solids in ppm

- Sulfate-> Amount of Sulfates dissolved in mg/L
- Conductivity-> Electrical conductivity of water in $\mu\text{S}/\text{cm}$
- Organic_carbon-> Amount of organic carbon in ppm
- Trihalomethanes-> Amount of Trihalomethanes in $\mu\text{g}/\text{L}$
- Turbidity-> Measure of light emitting property of water in NTU (Nephelometric Turbidity Units)
- Potability-> Indicates if water is safe for human consumption

In [3]: `data.describe()`

Out[3]:

	ph	Hardness	Solids	Chloramines	Sulfate	Conductivity
count	2785.000000	3276.000000	3276.000000	3276.000000	2495.000000	3276.000000
mean	7.080795	196.369496	22014.092526	7.122277	333.775777	426.205111
std	1.594320	32.879761	8768.570828	1.583085	41.416840	80.824064
min	0.000000	47.432000	320.942611	0.352000	129.000000	181.483754
25%	6.093092	176.850538	15666.690297	6.127421	307.699498	365.734414
50%	7.036752	196.967627	20927.833607	7.130299	333.073546	421.884968
75%	8.062066	216.667456	27332.762127	8.114887	359.950170	481.792304
max	14.000000	323.124000	61227.196008	13.127000	481.030642	753.342620

In [4]: `data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3276 entries, 0 to 3275
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ph                    2785 non-null   float64
1   Hardness              3276 non-null   float64
2   Solids                3276 non-null   float64
3   Chloramines           3276 non-null   float64
4   Sulfate               2495 non-null   float64
5   Conductivity          3276 non-null   float64
6   Organic_carbon        3276 non-null   float64
7   Trihalomethanes       3114 non-null   float64
8   Turbidity             3276 non-null   float64
9   Potability            3276 non-null   int64
dtypes: float64(9), int64(1)
memory usage: 256.1 KB
```


In [5]: `print('There are {} data points and {} features in the data'.format(data.shape[0], data.shape[1]))`

There are 3276 data points and 10 features in the data

Null Values

```
In [6]: sns.heatmap(data.isnull(),yticklabels=False,cbar=False,cmap='viridis')
```

Out[6]: <AxesSubplot:>



```
In [7]: for i in data.columns:
        if data[i].isnull().sum()>0:
            print("There are {} null values in {} column".format(data[i].isnull().sum(), i))
```

There are 491 null values in ph column

There are 781 null values in Sulfate column

There are 162 null values in Trihalomethanes column

Handling Null Values

PH

```
In [8]: data['ph'].describe()
```

```
Out[8]: count    2785.000000
        mean       7.080795
        std       1.594320
        min       0.000000
        25%       6.093092
        50%       7.036752
        75%       8.062066
        max      14.000000
        Name: ph, dtype: float64
```


Filling the missing values by mean

```
In [9]: data['ph_mean']=data['ph'].fillna(data['ph'].mean())
```

```
In [10]: data['ph_mean'].isnull().sum()
```

Out[10]: 0

```
In [11]: fig = plt.figure()
ax = fig.add_subplot(111)
data['ph'].plot(kind='kde', ax=ax)
data.ph_mean.plot(kind='kde', ax=ax, color='red')
lines, labels = ax.get_legend_handles_labels()
ax.legend(lines, labels, loc='best')
plt.show()
```




The distribution is not uniform

Filling the data with random values

```
In [12]: def impute_nan(df, variable):
df[variable+"_random"] = df[variable]
##It will have the random sample to fill the na
random_sample = df[variable].dropna().sample(df[variable].isnull().sum(), r
##pandas need to have same index in order to merge the dataset
random_sample.index = df[df[variable].isnull()].index
df.loc[df[variable].isnull(), variable+'_random'] = random_sample
```

```
In [13]: impute_nan(data, "ph")
```


```
In [14]: fig = plt.figure()
ax = fig.add_subplot(111)
data['ph'].plot(kind='kde', ax=ax)
data.ph_random.plot(kind='kde', ax=ax, color='green')
data.ph_mean.plot(kind='kde', ax=ax, color='red')
lines, labels = ax.get_legend_handles_labels()
ax.legend(lines, labels, loc='best')
plt.show()
```



Uniform distribution with random initialization

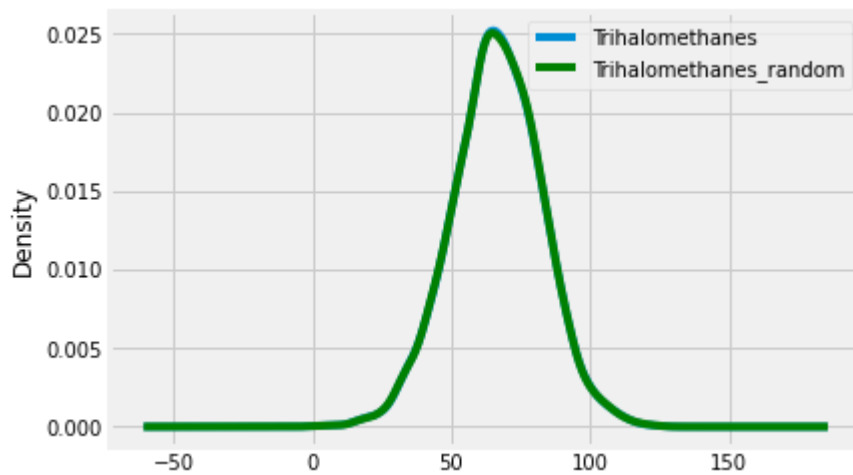
```
In [15]: impute_nan(data, "Sulfate")
```

```
In [16]: fig = plt.figure()
ax = fig.add_subplot(111)
data['Sulfate'].plot(kind='kde', ax=ax)
data["Sulfate_random"].plot(kind='kde', ax=ax, color='green')
#data.ph_mean.plot(kind='kde', ax=ax, color='red')
lines, labels = ax.get_legend_handles_labels()
ax.legend(lines, labels, loc='best')
plt.show()
```



```
In [17]: impute_nan(data, "Trihalomethanes")
```

```
In [18]: fig = plt.figure()
ax = fig.add_subplot(111)
data['Trihalomethanes'].plot(kind='kde', ax=ax)
data.Trihalomethanes_random.plot(kind='kde', ax=ax, color='green')
lines, labels = ax.get_legend_handles_labels()
ax.legend(lines, labels, loc='best')
plt.show()
```




```
In [19]: data=data.drop(['ph','Sulfate','Trihalomethanes','ph_mean'],axis=1)
```

```
In [20]: data.isnull().sum()
```

```
Out[20]: Hardness          0
         Solids            0
         Chloramines       0
         Conductivity      0
         Organic_carbon    0
         Turbidity         0
         Potability        0
         ph_random         0
         Sulfate_random    0
         Trihalomethanes_random  0
         dtype: int64
```

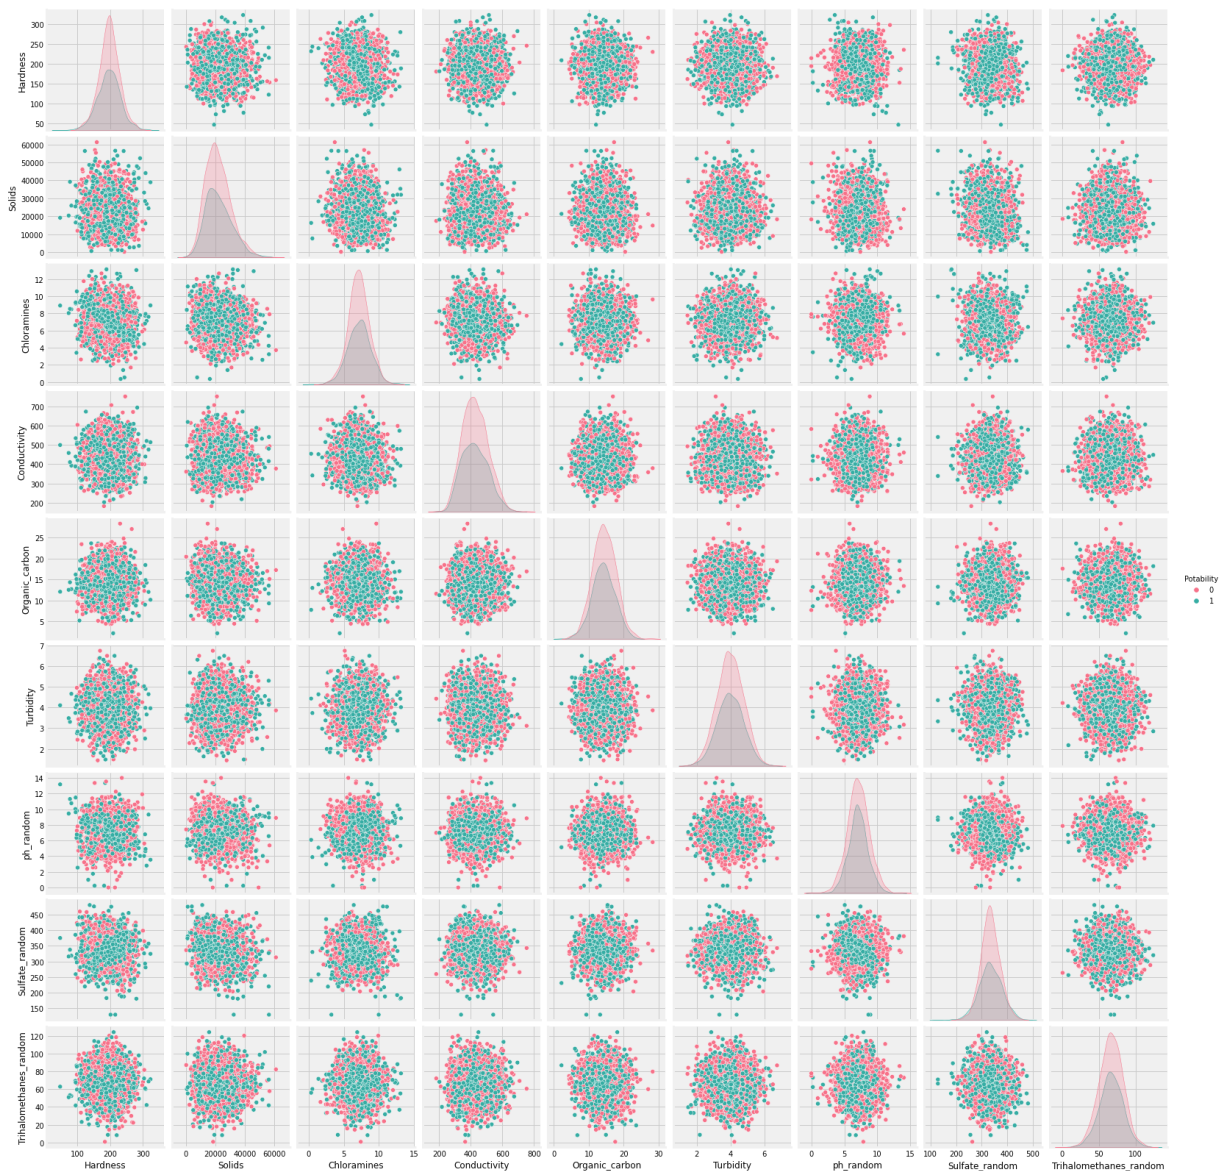
Check for Correlation

```
In [21]: plt.figure(figsize=(20, 17))
         matrix = np.triu(data.corr())
         sns.heatmap(data.corr(), annot=True,linewidth=.8, mask=matrix, cmap="rocket")
```




There are no correlated columns presebt in the data

```
In [22]: sns.pairplot(data, hue="Potability", palette="husl");
```



```
In [23]: def distributionPlot(data):
    """
    Creates distribution plot.
    """
    fig = plt.figure(figsize=(20, 20))
    for i in tqdm_notebook(range(0, len(data.columns))):
        fig.add_subplot(np.ceil(len(data.columns)/3), 3, i+1)
        sns.distplot(
            data.iloc[:, i], color="lightcoral", rug=True)
        fig.tight_layout(pad=3.0)
    plot_data = data.drop(['Potability'], axis = 1)
    distributionPlot(plot_data)
```

0% | 0/9 [00:00<?, ?it/s]




Hardness

```
In [24]: data['Hardness'].describe()
```


```
Out[24]: count    3276.000000
         mean      196.369496
         std       32.879761
         min       47.432000
         25%      176.850538
         50%      196.967627
         75%      216.667456
         max       323.124000
         Name: Hardness, dtype: float64
```

```
In [25]: plt.figure(figsize = (16, 7))
         sns.distplot(data['Hardness'])
```

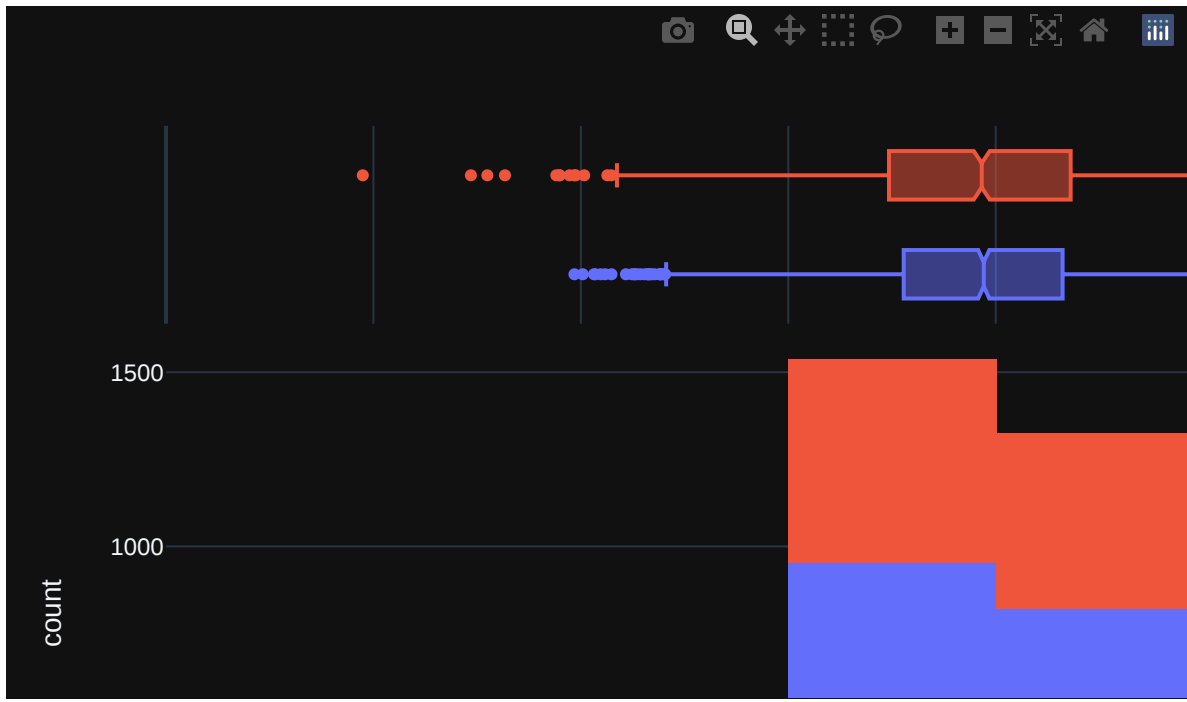
```
plt.title('Distribution Plot of Hardness\n', fontsize = 20)
plt.show()
```



```
In [26]: # basic scatter plot
fig = px.scatter(data, range(data['Hardness'].count()), sorted(data['Hardness'],
    color=data['Potability'],
    labels={
        'x': "Count",
        'y': "Hardness",
        'color': 'Potability'
    }, template = 'plotly_dark')
fig.update_layout(title='Hardness wrt Potability')
fig.show()
```



```
In [27]: px.histogram(data_frame = data, x = 'Hardness', nbins = 10, color = 'Potabil  
         template = 'plotly_dark')
```




Solids

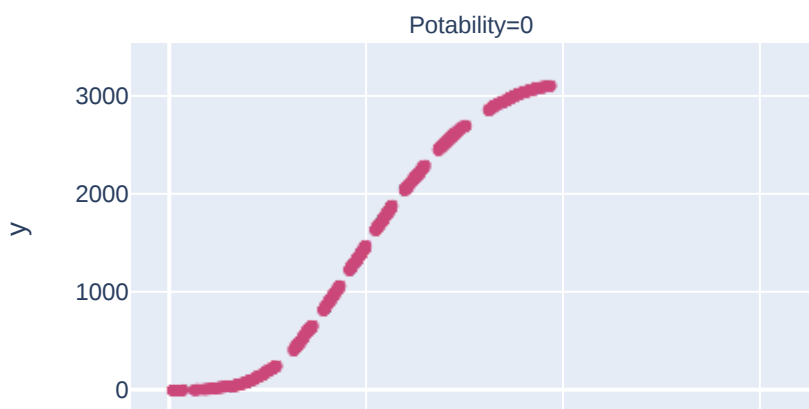
```
In [28]: data['Solids'].describe()
```

```
Out[28]: count      3276.000000
mean      22014.092526
std       8768.570828
min        320.942611
25%      15666.690297
50%      20927.833607
75%      27332.762127
max       61227.196008
Name: Solids, dtype: float64
```


```
In [29]: plt.figure(figsize = (16, 7))
sns.distplot(data['Solids'])
plt.title('Distribution Plot of Solids\n', fontsize = 20)
plt.show()
```




```
In [30]: fig = px.scatter(data, sorted(data["Solids"]), range(data["Solids"].count())
          facet_row="Potability")
          fig.show()
```



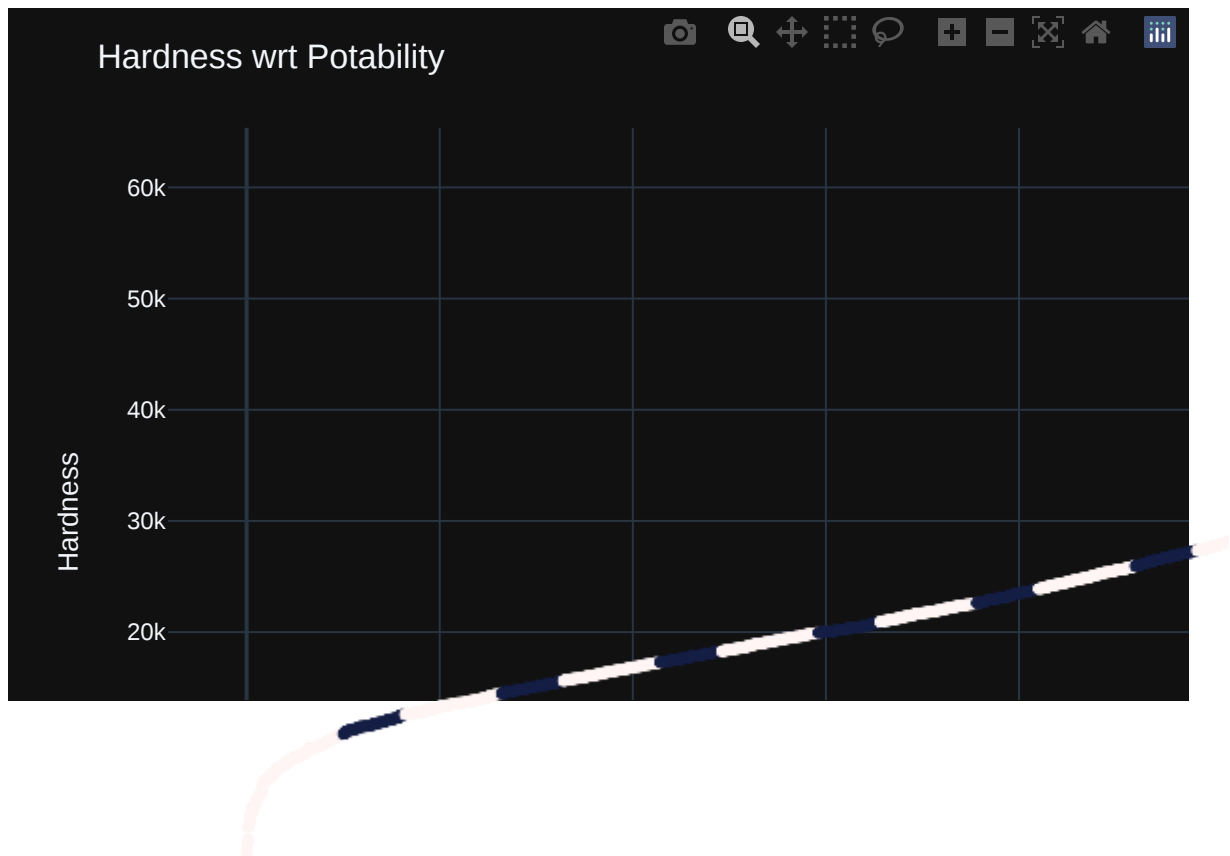
Po



```
In [31]: px.histogram(data_frame = data, x = 'Solids', nbins = 10, color = 'Potabilit
          template = 'plotly_dark')
```



```
In [32]: # basic scatter plot
fig = px.scatter(data, range(data['Solids'].count()), sorted(data['Solids']),
                 color=data['Potability'],
                 labels={
                     'x': "Count",
                     'y': "Hardness",
                     'color': 'Potability'
                 },
                 color_continuous_scale=px.colors.sequential.tempo,
                 template = 'plotly_dark')
fig.update_layout(title='Hardness wrt Potability')
fig.show()
```



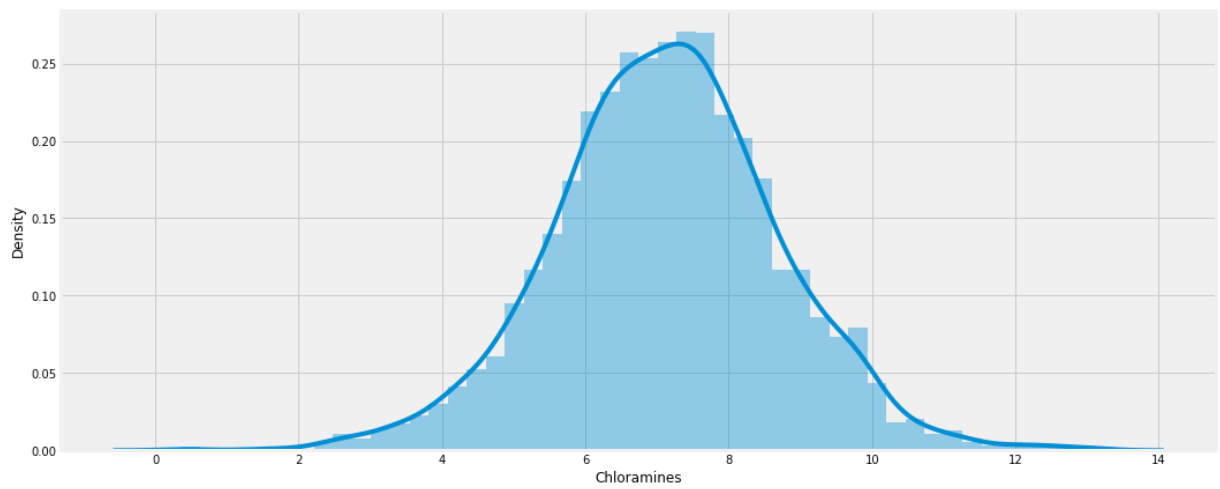
Chloramines

```
In [33]: data['Chloramines'].describe()
```

```
Out[33]: count    3276.000000
mean         7.122277
std          1.583085
min          0.352000
25%          6.127421
50%          7.130299
75%          8.114887
max          13.127000
Name: Chloramines, dtype: float64
```

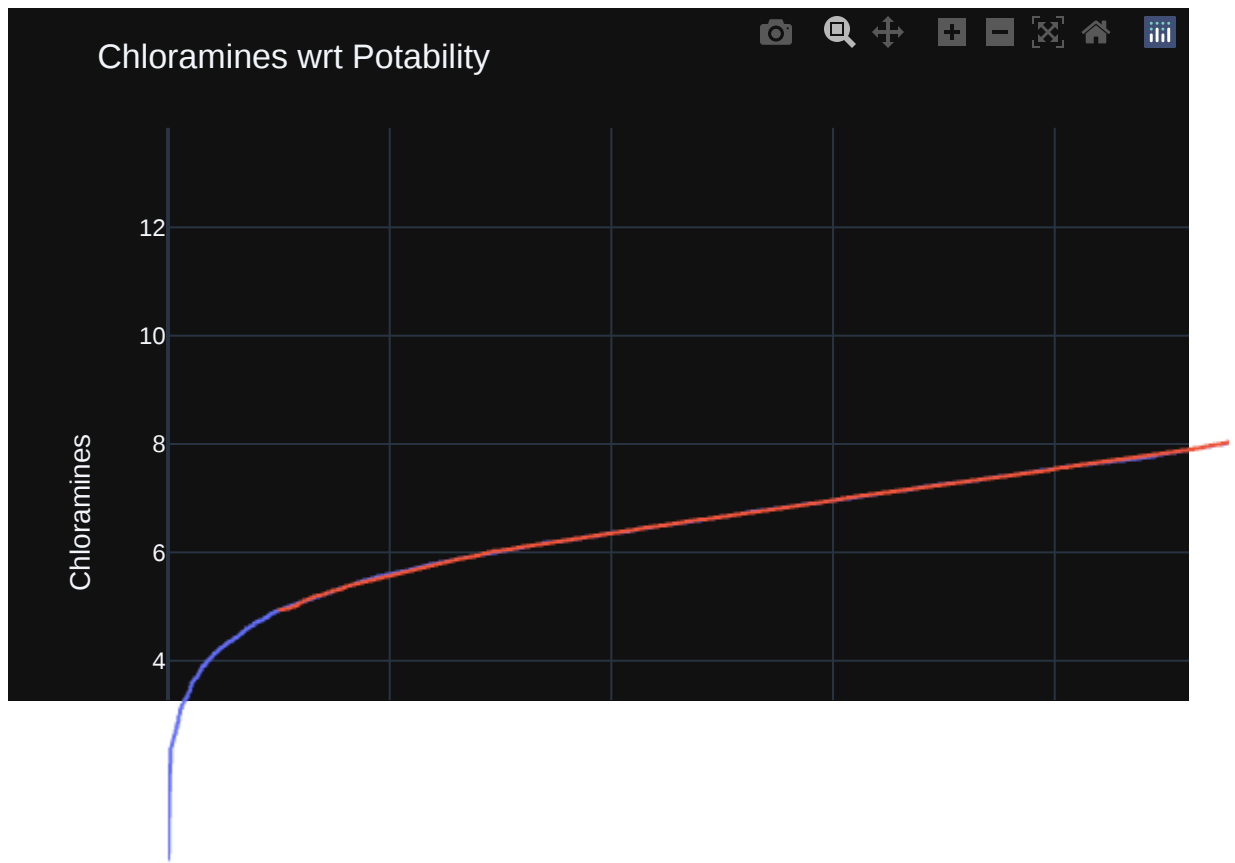
```
In [34]: plt.figure(figsize = (16, 7))
sns.distplot(data['Chloramines'])
plt.title('Distribution Plot of Chloramines\n', fontsize = 20)
plt.show()
```

Distribution Plot of Chloramines

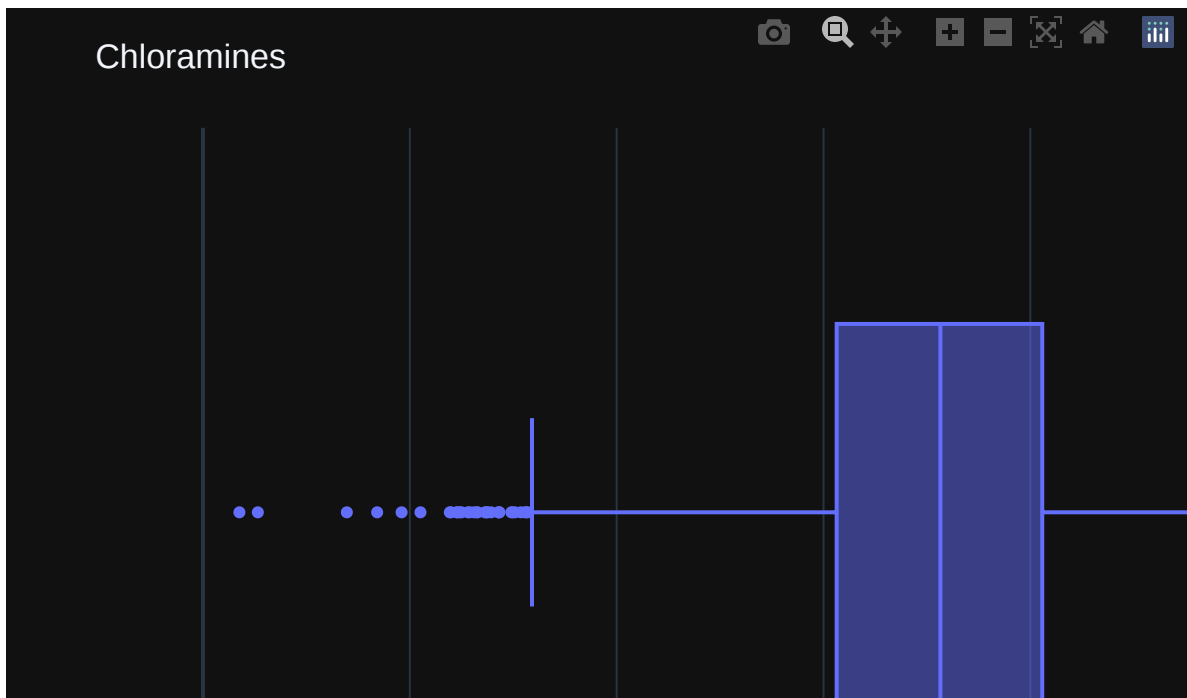


```
In [35]: fig = px.line(x=range(data['Chloramines'].count()), y=sorted(data['Chloramines']),
                        'x': "Count",
                        'y': "Chloramines",
                        'color': 'Potability'

                        }, template = 'plotly_dark')
fig.update_layout(title='Chloramines wrt Potability')
fig.show()
```

```
In [36]: fig = px.box(x = 'Chloramines', data_frame = data, template = 'plotly_dark')
fig.update_layout(title='Chloramines')
fig.show()
```




Conductivity

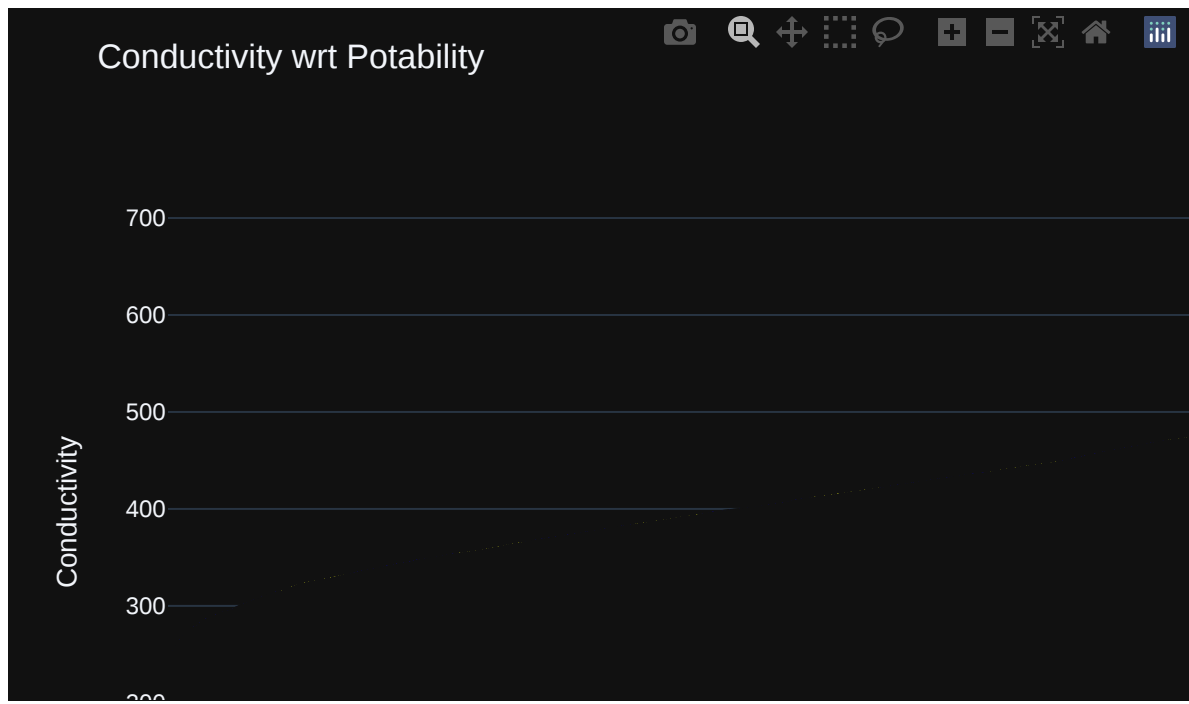
```
In [37]: data["Conductivity"].describe()
```

```
Out[37]: count    3276.000000
mean      426.205111
std       80.824064
min       181.483754
25%       365.734414
50%       421.884968
75%       481.792304
max       753.342620
Name: Conductivity, dtype: float64
```


```
In [38]: plt.figure(figsize = (16, 7))
sns.distplot(data['Conductivity'])
plt.title('Distribution Plot of Conductivity\n', fontsize = 20)
plt.show()
```



```
In [39]: fig = px.bar(data, x=range(data['Conductivity'].count()),
                    y=sorted(data['Conductivity']), labels={
                        'x': "Count",
                        'y': "Conductivity",
                        'color': 'Potability'
                    },
                    color=data['Potability'],
                    template = 'plotly_dark')
fig.update_layout(title='Conductivity wrt Potability')
fig.show()
```



```
In [40]: group_labels = ['distplot'] # name of the dataset  
  
fig = ff.create_distplot([data['Conductivity']], group_labels)  
fig.show()
```




Organic_carbon

```
In [41]: data['Organic_carbon'].describe()
```

```
Out[41]: count    3276.000000
         mean      14.284970
         std       3.308162
         min       2.200000
         25%      12.065801
         50%      14.218338
         75%      16.557652
         max      28.300000
         Name: Organic_carbon, dtype: float64
```

```
In [42]: group_labels = ['Organic_carbon'] # name of the dataset


         fig = ff.create_distplot([data['Organic_carbon']], group_labels)
         fig.show()
```




```
In [43]: dt_5=data[data['Organic_carbon']<5]
dt_5_10=data[(data['Organic_carbon']>5)&(data['Organic_carbon']<10)]
dt_10_15=data[(data['Organic_carbon']>10)&(data['Organic_carbon']<15)]
dt_15_20=data[(data['Organic_carbon']>15)&(data['Organic_carbon']<20)]
dt_20_25=data[(data['Organic_carbon']>20)&(data['Organic_carbon']<25)]
dt_25=data[(data['Organic_carbon']>25)]

x_Age = ['5', '5-10', '10-15', '15-20', '25+']
y_Age = [len(dt_5.values), len(dt_5_10.values), len(dt_10_15.values), len(dt_15_20.values), len(dt_20_25.values)]

px.bar(data_frame = data, x = x_Age, y = y_Age, color = x_Age, template = 'px',
       title = 'Number of passengers per Age group')
```



```
In [44]: sns.catplot(x = 'Organic_carbon', y = 'Organic_carbon', hue = 'Potability',  
                    height = 5, aspect = 2)  
plt.show()
```




Turbidity

```
In [45]: data['Turbidity'].describe()
```

```
Out[45]: count    3276.000000  
mean         3.966786  
std          0.780382  
min          1.450000  
25%          3.439711  
50%          3.955028  
75%          4.500320  
max          6.739000  
Name: Turbidity, dtype: float64
```

```
In [46]: group_labels = ['Turbidity'] # name of the dataset  
  
fig = ff.create_distplot([data['Turbidity']], group_labels)  
fig.show()
```




```
In [47]: data['turbid_class']=data['Turbidity'].astype(int)
```



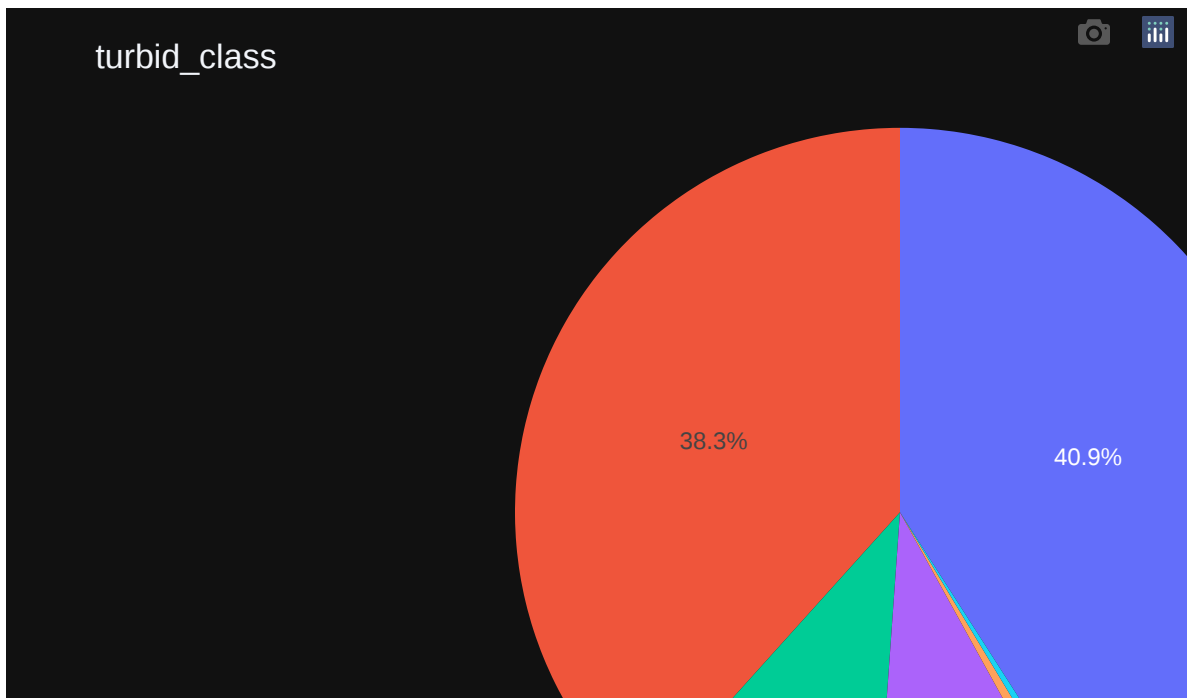
```
In [48]: data['turbid_class'].unique()
```

```
Out[48]: array([2, 4, 3, 5, 6, 1])
```

```
In [49]: px.scatter(data_frame = data, x = 'Turbidity', y = 'turbid_class', color = '')
```



```
In [50]: fig = px.pie(data,
                      values=data['turbid_class'].value_counts(),
                      names=data['turbid_class'].value_counts().keys(),
                      )
fig.update_layout(
    title='turbid_class',
    template = 'plotly_dark'
)
fig.show()
```



```
In [51]: data=data.drop(['turbid_class'],axis=1)
```


ph_random

```
In [52]: data['ph_random'].describe()
```

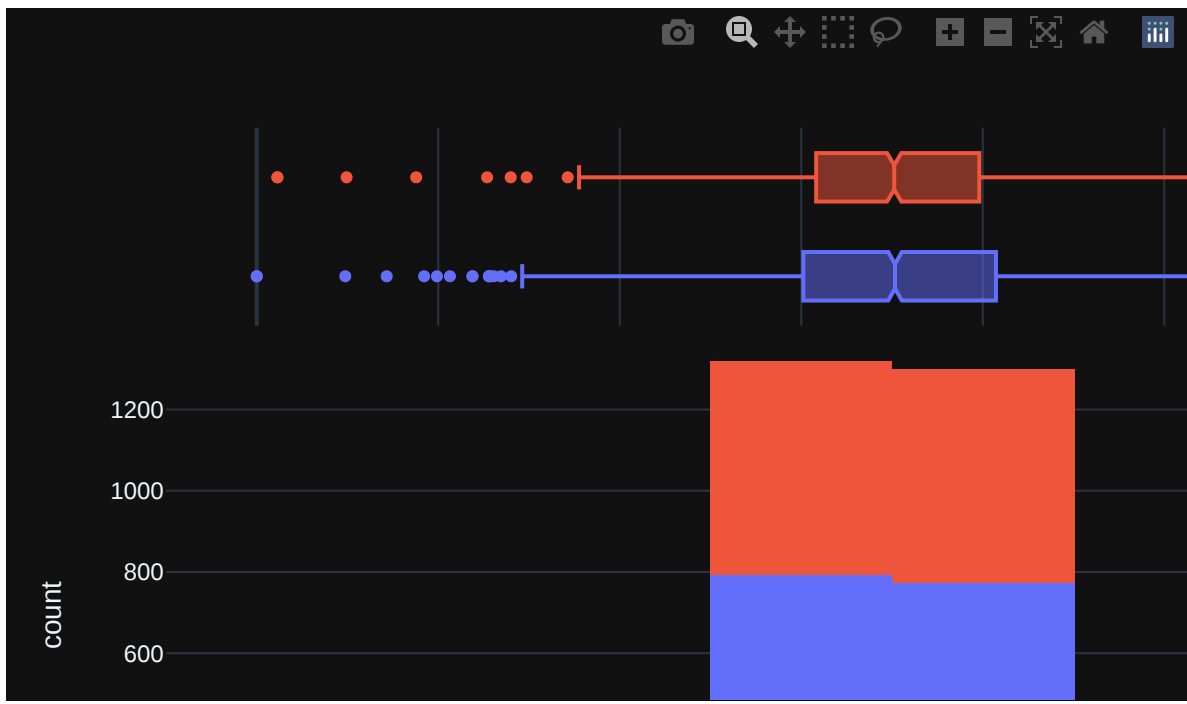
```
Out[52]: count    3276.000000
mean         7.071639
std          1.607991
min           0.000000
25%          6.081460
50%          7.029490
75%          8.063147
max          14.000000
Name: ph_random, dtype: float64
```

```
In [53]: group_labels = ['ph_random'] # name of the dataset


fig = ff.create_distplot([data['ph_random']], group_labels)
fig.show()
```



```
In [54]: px.histogram(data_frame = data, x = 'ph_random', nbins = 10, color = 'Potabi  
          template = 'plotly_dark')
```



```
In [55]: fig = px.scatter(data, sorted(data["ph_random"]), range(data["ph_random"].count(), data["ph_random"].max()),
                        facet_row="Potability")
fig.show()
```




Sulfate_random

```
In [56]: data['Sulfate_random'].describe()
```


```
Out[56]: count    3276.000000
         mean      333.430954
         std       41.026947
         min       129.000000
         25%       307.523159
         50%       332.879578
         75%       359.710517
         max       481.030642
         Name: Sulfate_random, dtype: float64
```

```
In [57]: group_labels = ['distplot'] # name of the dataset

         fig = ff.create_distplot([data['Sulfate_random']], group_labels)
         fig.show()
```



```
In [58]: sns.catplot(x = 'Sulfate_random', y = 'Sulfate_random', hue = 'Potability',
                    height = 5, aspect = 2)
plt.show()
```




Trihalomethanes_random

```
In [59]: data['Trihalomethanes_random'].describe()
```


```
Out[59]: count    3276.000000  
mean       66.419200  
std        16.184832  
min         0.738000  
25%        55.861675  
50%        66.639068  
75%        77.384166  
max       124.000000  
Name: Trihalomethanes_random, dtype: float64
```

```
In [60]: group_labels = ['Trihalomethanes_random'] # name of the dataset  
  
fig = ff.create_distplot([data['Trihalomethanes_random']], group_labels)  
fig.show()
```




```
In [61]: fig = px.box(x = 'Trihalomethanes_random', data_frame = data, template = 'pl  
fig.update layout(title='Trihalomethanes_random')
```

```
fig.show()
```



```
In [62]: fig = px.line(x=range(data['Trihalomethanes_random'].count()), y=sorted(data
                                'x': "Count",
                                'y': "Trihalomethanes",
                                'color': 'Potability'
                                }, template = 'plotly_dark')
fig.update_layout(title='Trihalomethane wrt Potability')
fig.show()
```





Potability


```
In [63]: data['Potability'].describe()
```

```
Out[63]: count    3276.000000
         mean      0.390110
         std       0.487849
         min       0.000000
         25%       0.000000
         50%       0.000000
         75%       1.000000
         max       1.000000
         Name: Potability, dtype: float64
```

```
In [64]: px.histogram(data_frame = data, x = 'Potability', color = 'Potability', margin = 10,
                       template = 'plotly_dark')
```



```
In [65]: fig = px.pie(data,
                    values=data['Potability'].value_counts(),
                    names=data['Potability'].value_counts().keys(),
                    )
fig.update_layout(
    title='Potability',
    template = 'plotly_dark'
)
fig.show()
```



Data Preprocessing

```
In [66]: from sklearn.preprocessing import StandardScaler  
from sklearn.model_selection import train_test_split
```

```
In [67]: X=data.drop(['Potability'],axis=1)  
y=data['Potability']
```

Since the data is not in a uniform shape, we scale the data using standard scalar

```
In [68]: scaler = StandardScaler()  
x=scaler.fit_transform(X)
```

```
In [69]: # split the data to train and test set  
x_train,x_test,y_train,y_test = train_test_split(x,y,train_size=0.85,random_  
  
print("training data shape:-{} labels{}".format(x_train.shape,y_train.shape)  
print("testing data shape:-{} labels{}".format(x_test.shape,y_test.shape))
```

```
training data shape:-(2784, 9) labels(2784,)
testing data shape:-(492, 9) labels(492,)
```


Modeling

Logistic Regression

```
In [70]: from sklearn.linear_model import LogisticRegression
log = LogisticRegression(random_state=0).fit(x_train, y_train)
log.score(x_test, y_test)
```

```
Out[70]: 0.6219512195121951
```

```
In [71]: # Confusion matrix
from sklearn.metrics import confusion_matrix
# Make Predictions
pred1=log.predict(np.array(x_test))
plt.title("Confusion Matrix testing data")
sns.heatmap(confusion_matrix(y_test,pred1),annot=True,cbar=False)
plt.legend()
plt.show()
```



K Nearest Neighbours

```
In [72]: from sklearn.neighbors import KNeighborsClassifier
```


```
In [73]: knn = KNeighborsClassifier(n_neighbors=2)
# Train the model using the training sets
knn.fit(x_train,y_train)

#Predict Output
predicted= knn.predict(x_test) # 0:Overcast, 2:Mild
```

```
In [74]: # Confusion matrix
```

```
Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js fusion_matrix
```

```
# Make Predictions
pred1=knn.predict(np.array(x_test))
plt.title("Confusion Matrix testing data")
sns.heatmap(confusion_matrix(y_test,pred1),annot=True,cbar=False)
plt.legend()
plt.show()
```



SVM


```
In [75]: from sklearn import svm
from sklearn.metrics import accuracy_score
```

```
In [76]: svmc = svm.SVC()
svmc.fit(x_train, y_train)

y_pred = svmc.predict(x_test)
print(accuracy_score(y_test,y_pred))
```

0.6808943089430894

```
In [77]: # Confusion matrix
from sklearn.metrics import confusion_matrix
# Make Predictions
pred1=svmc.predict(np.array(x_test))
plt.title("Confusion Matrix testing data")
sns.heatmap(confusion_matrix(y_test,pred1),annot=True,cbar=False)
plt.legend()
plt.show()
```



Decision Tree


```
In [78]: from sklearn import tree
        from sklearn.metrics import accuracy_score
```

```
In [79]: tre = tree.DecisionTreeClassifier()
        tre = tre.fit(x_train, y_train)

        y_pred = tre.predict(x_test)
        print(accuracy_score(y_test, y_pred))
```

0.5487804878048781

```
In [80]: # Confusion matrix
        from sklearn.metrics import confusion_matrix
        # Make Predictions
        pred1=tre.predict(np.array(x_test))
        plt.title("Confusion Matrix testing data")
        sns.heatmap(confusion_matrix(y_test, pred1), annot=True, cbar=False)
        plt.legend()
        plt.show()
```



Random Forest

```
In [81]: from sklearn.ensemble import RandomForestClassifier
        from sklearn.metrics import accuracy_score
```


```
In [82]: # create the model
        model_rf = RandomForestClassifier(n_estimators=500, oob_score=True, random_s

        # fitting the model
        model_rf=model_rf.fit(x_train, y_train)

        y_pred = model_rf.predict(x_test)
        print(accuracy_score(y_test,y_pred))
```

0.6788617886178862

```
In [83]: # Confusion matrix
        from sklearn.metrics import confusion_matrix
        # Make Predictions
        pred1=model_rf.predict(np.array(x_test))
        plt.title("Confusion Matrix testing data")
        sns.heatmap(confusion_matrix(y_test,pred1),annot=True,cbar=False)
        plt.legend()
        plt.show()
```



XG Boost

```
In [84]: from xgboost import XGBClassifier
        from sklearn.metrics import r2_score

        xgb = XGBClassifier(colsample_bylevel= 0.9,
                             colsample_bytree = 0.8,
                             gamma=0.99,
                             max_depth= 5,
                             min_child_weight= 1,
                             n_estimators= 8,
```

```

        nthread= 5,
        random_state= 0,
    )
xgb.fit(x_train,y_train)

```

[14:31:40] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

```

Out[84]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=0.9,
        colsample_bynode=1, colsample_bytree=0.8, gamma=0.99, gpu_id=-1,
        importance_type='gain', interaction_constraints='',
        learning_rate=0.300000012, max_delta_step=0, max_depth=5,
        min_child_weight=1, missing=nan, monotone_constraints=(),
        n_estimators=8, n_jobs=5, nthread=5, num_parallel_tree=1,
        random_state=0, reg_alpha=0, reg_lambda=1, scale_pos_weight=1,
        subsample=1, tree_method='exact', validate_parameters=1,
        verbosity=None)

```

```

In [85]: print('Accuracy of XGBoost classifier on training set: {:.2f}'
        .format(xgb.score(x_train, y_train)))
        print('Accuracy of XGBoost classifier on test set: {:.2f}'
        .format(xgb.score(x_test, y_test)))

```


Accuracy of XGBoost classifier on training set: 0.72
 Accuracy of XGBoost classifier on test set: 0.63

```

In [86]: from sklearn.metrics import confusion_matrix

conf_matrix = confusion_matrix(y_true=y_test, y_pred=y_pred)
plt.figure(figsize = (15, 8))
sns.set(font_scale=1.4) # for label size
sns.heatmap(conf_matrix, annot=True, annot_kws={"size": 16}, cbar=False, line
plt.title("Test Confusion Matrix")
plt.xlabel("Predicted class")
plt.ylabel("Actual class")
plt.savefig('conf_test.png')
plt.show()

```

SVM tuned

```
In [87]: from sklearn.svm import SVC
from sklearn.model_selection import GridSearchCV
svc=SVC()
param_grid={'C':[1.2,1.5,2.2,3.5,3.2,4.1], 'kernel':['linear', 'poly', 'rbf'],
gridsearch=GridSearchCV(svc,param_grid=param_grid,n_jobs=-1,verbose=4,cv=3)
gridsearch.fit(x_train,y_train)
```

Fitting 3 folds for each of 240 candidates, totalling 720 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=-1)]: Done 17 tasks      | elapsed: 2.9s
[Parallel(n_jobs=-1)]: Done 90 tasks      | elapsed: 7.5s
[Parallel(n_jobs=-1)]: Done 213 tasks     | elapsed: 15.3s
[Parallel(n_jobs=-1)]: Done 384 tasks     | elapsed: 26.8s
[Parallel(n_jobs=-1)]: Done 605 tasks     | elapsed: 42.5s
[Parallel(n_jobs=-1)]: Done 720 out of 720 | elapsed: 51.0s finished
```

```
Out[87]: GridSearchCV(cv=3, estimator=SVC(), n_jobs=-1,
                    param_grid={'C': [1.2, 1.5, 2.2, 3.5, 3.2, 4.1],
                                'degree': [1, 2, 4, 8, 10], 'gamma': ['scale', 'auto'],
                                'kernel': ['linear', 'poly', 'rbf', 'sigmoid']}),
                    verbose=4)
```

```
In [88]: y_pred=gridsearch.predict(x_test)
from sklearn.metrics import confusion_matrix

conf_matrix = confusion_matrix(y_true=y_test, y_pred=y_pred)
plt.figure(figsize = (15, 8))
sns.set(font_scale=1.4) # for label size
sns.heatmap(conf_matrix, annot=True, annot_kws={"size": 16},cbar=False, line
plt.title("Test Confusion Matrix")
plt.xlabel("Predicted class")
plt.ylabel("Actual class")
```

```
plt.savefig('conf_test.png')  
plt.show()
```

